

Robotics

Austin, Oliver, Jeffrey

Table of Contents

Introduction

Robotics Software (RoboSuite, RoboVerse, LeRobot)

Generalist Policies (OpenVLA, Pi0)

Robotics in 2025

Eric Jang - “all roads lead to robotics”

Classical to Robot Learning

Data as a fossil fuel

Democratized Hardware

Rapid industry interest



Robosuite (ARISE)

2017, SVL, now monitored by NVIDIA Gear, UT Austin PRL, and SVL.

(GDM) MuJoCo-based simulator

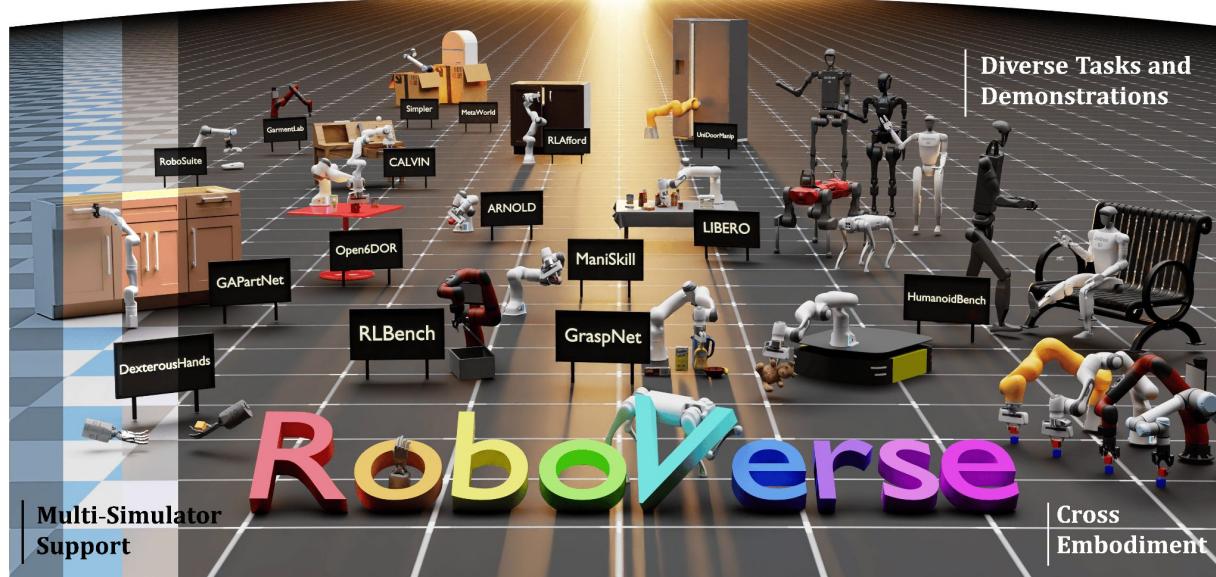
Most popular choice in academia

Standardized benchmarks, policy training,
imitation learning



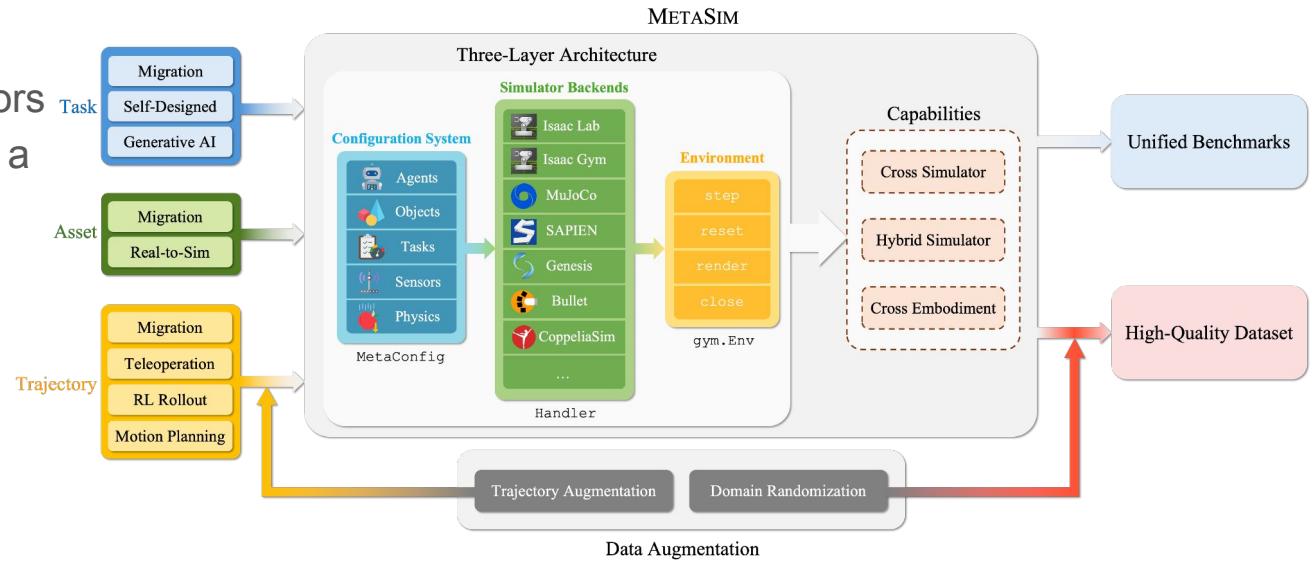
Roboverse (BAIR)

April 26, 2025 by BAIR and Peking et. al



RoboVerse - MetaSIM

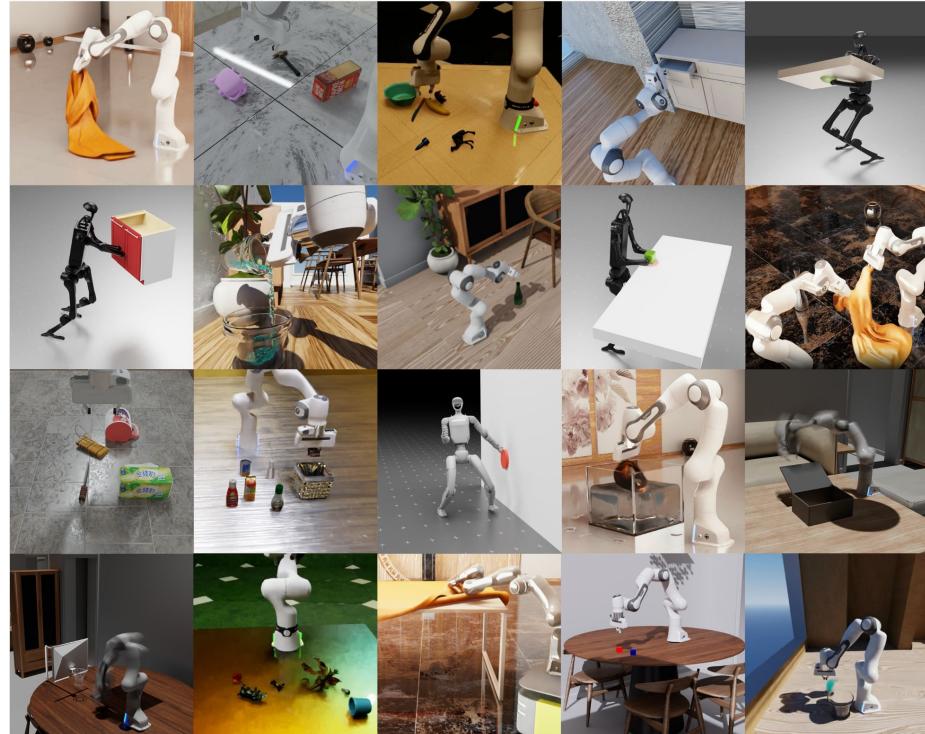
Unified simulation platform combining existing simulators and rendering engines into a single framework.



RoboVerse - Dataset

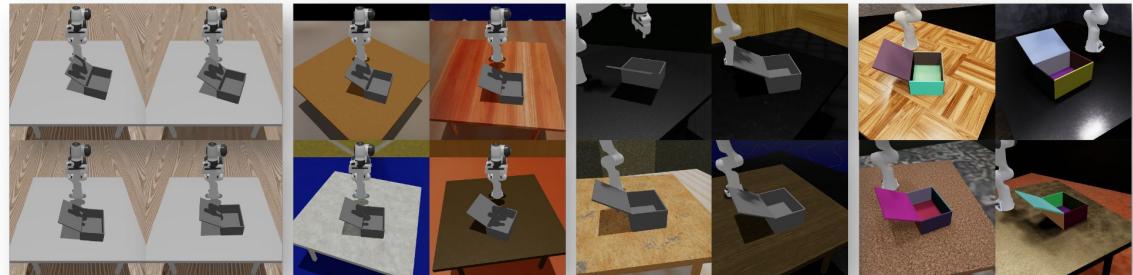
Large-scale, high-quality dataset.

Upon training VLA, shown to generalized to unseen tasks and new environments.



RoboVerse - Benchmark

Standardized benchmark for both imitation learning and reinforcement learning.



(a) Task Space

(b) Environment

(c) Camera

(d) Material & Lighting

LeRobot (Hugging Face)

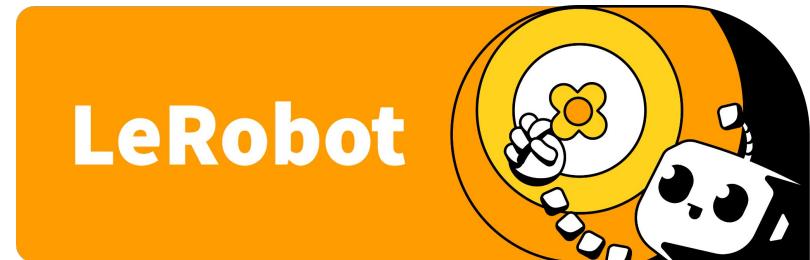
PyTorch based pipeline aimed to integrate the robotics stack into a single hub

LeRobot Dataset

Very active open community via HF

Everything is pushed to a central hub

Open source for hardware as well



 **AI & ML interests**
State-of-the-art Machine Learning for real-world robotics

 **Recent Activity**

-  **jadechoghari** updated a dataset 1 day ago
[lerobot/behavior1k-task0009](#)
-  **jadechoghari** updated a dataset 1 day ago
[lerobot/behavior1k-task0008](#)
-  **jadechoghari** updated a dataset 1 day ago
[lerobot/behavior1k-task0007](#)

[View all activity](#)

Limitations and Use Case

	Robosuite	Roboverse	LeRobot
Simulator	MuJoCo	Meta-Sim (all)	All, but need to port yourself
Maturity	Stable	Developing	Developing Fast
Task Scope	Table-top manipulation	Broad manipulation across simulations	Real-world + whatever sims you attach
Benchmarks	Common baselines	Large synthetic dataset	Hub-native dataset
Policy types	IL / RL	IL / RL	IL / RL + VLA friendly

Hardware Bottlenecks

Settled on humanoid form factor

Problems:

Dexterous Manipulation

Battery Life and Thermal management

GPU and VRAM Limitations

Cost, but there's Unitree

Data Foundry



References

<https://roboverseorg.github.io/>

<https://huggingface.co/spaces/lerobot/robot-learning-tutorial>

<https://github.com/huggingface/lerobot>

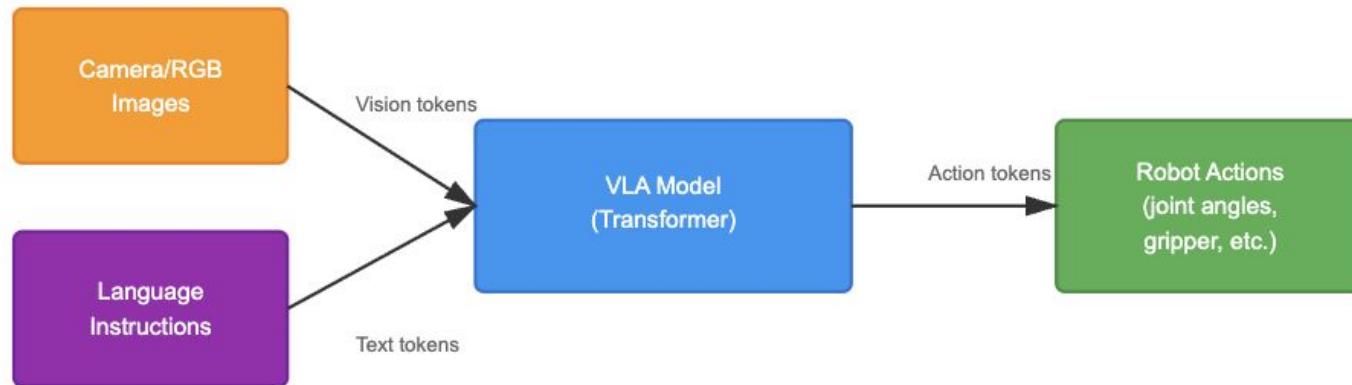
<https://robosuite.ai/>

<https://evjang.com/2024/03/03/all-roads-robots.html>

Generalist Policies

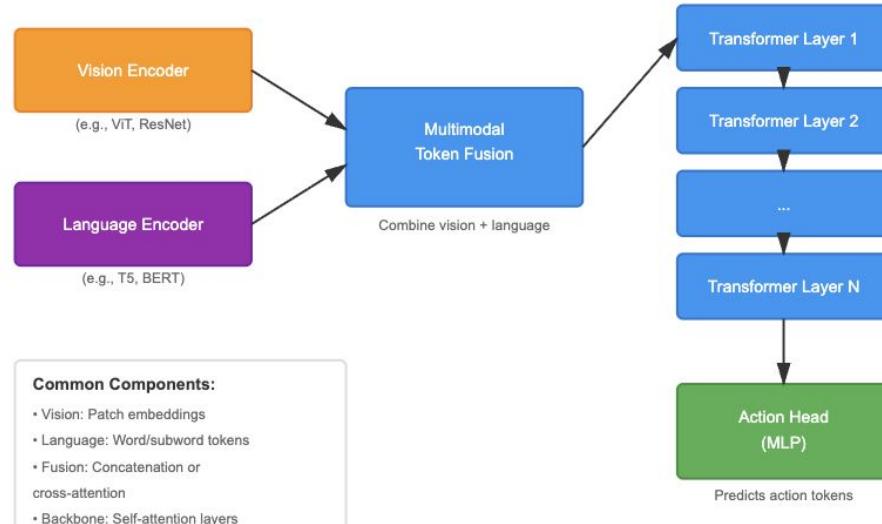
What is a VLA?

- Vision Language Action (VLA) Models build on VLMs by outputting action instead of text



VLA Architecture

- Combine image and text processing capabilities
- Example: Combine vision transformer and BERT tokenizer, fuse the tokens, pass to a transformer architecture.
- Transformer is trained to output contextualized representations
- MLP is trained to convert those into robot actions



Examples

Popular VLA Models:

RT-2 (Robotic Transformer 2): Uses vision-language model (PaLI-X) backbone, outputs discretized actions

OpenVLA: Open-source VLA based on LLaVA architecture, 7B parameters

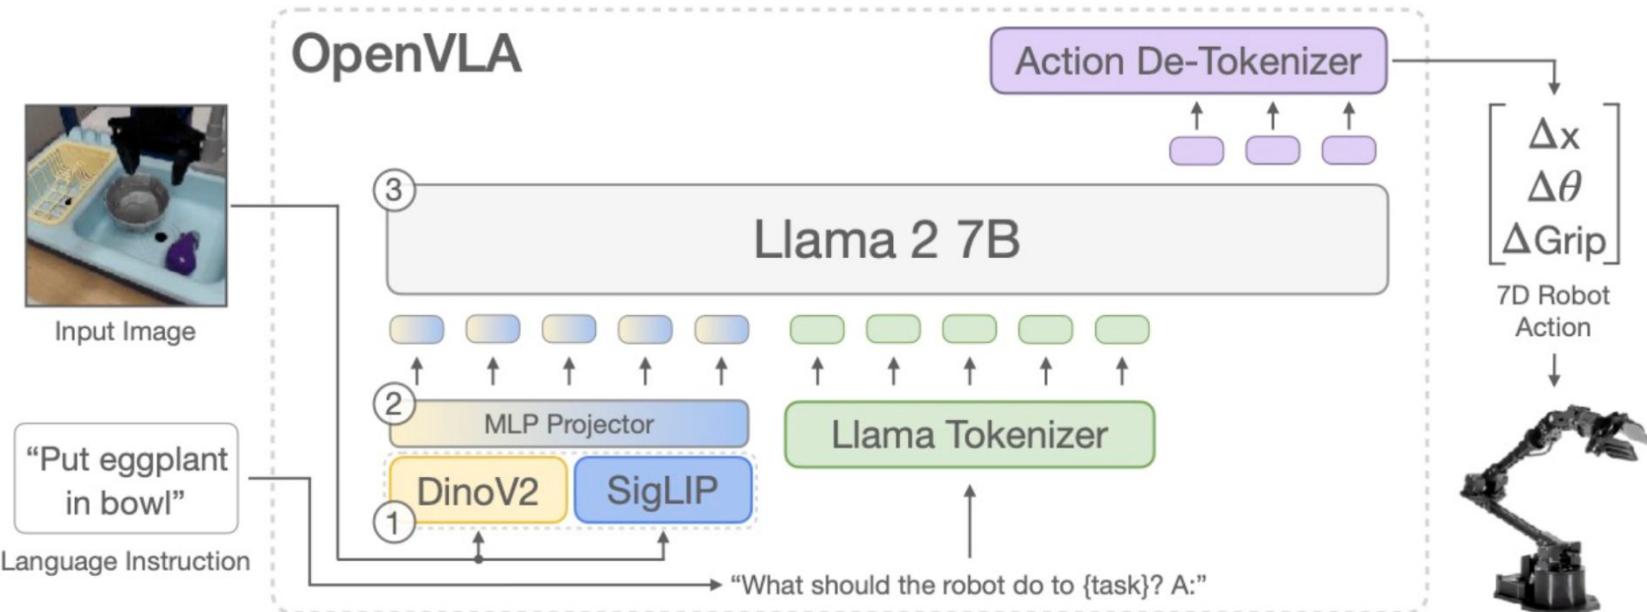
Octo: Generalist policy trained on 800k robot trajectories

PIVOT: Uses iterative planning with VLAs for complex tasks

OpenVLA

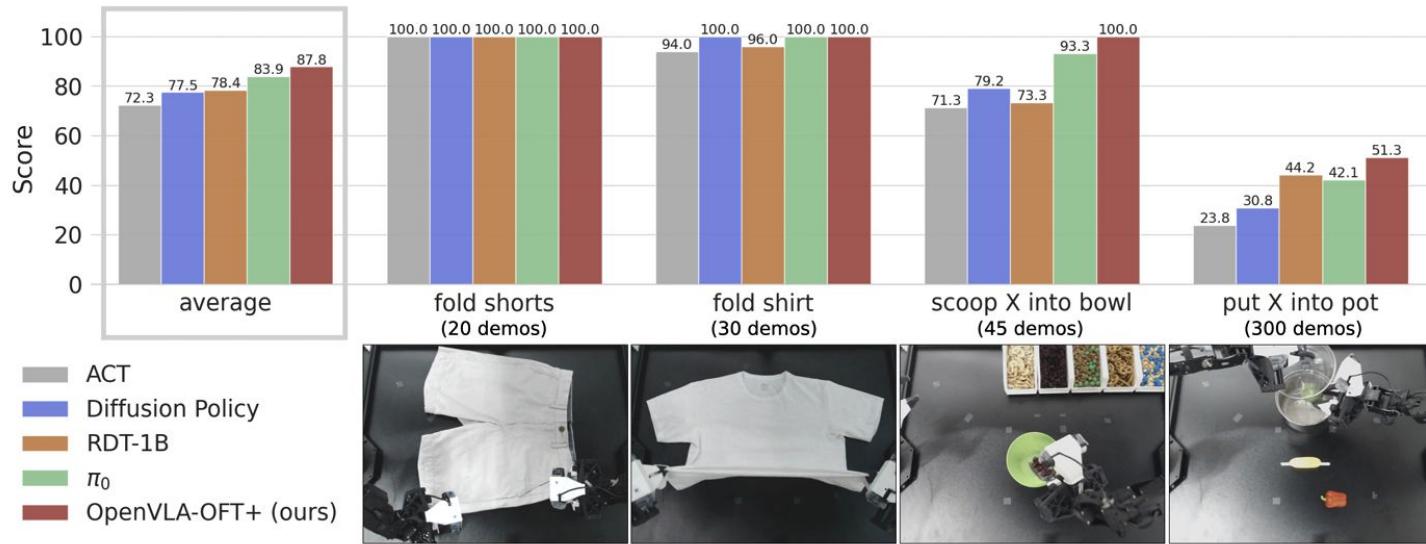
- First open-source VLA model.
- OpenVLA consists of four key components:
 - Fused visual encoder combining SigLIP and DINOv2 backbones that processes image inputs into patch embeddings
 - Built in projector maps these embeddings into the language model's input space
 - Llama-2-7B language model backbone that predicts tokenized output actions.
 - Discrete actions processed through softmax-based token prediction
- 256 discrete action tokens from the Llama vocabulary to represent robot control values with 8-bit resolution, covering a 7-DoF action space.

OpenVLA Architecture



OpenVLA-OFT

- Optimized for fine tuning
 - Parallelized instead of autoregressive decoding
 - Continuous outputs instead of discrete outputs
- Compares L1 Regression and Diffusion approaches
 - L1 is easier to calculate (requires less compute)
 - L1 converges faster
 - Diffusion can do really poorly when some training examples are incorrect, L1 does a better job smoothing out the noise



OpenVLA, a 7B-parameter model, outperforms the 55B-parameter RT-2-X model by 16.5% absolute success rate across 29 evaluation tasks on the WidowX and Google Robot embodiments

Fine-tuning on LIBERO

Checkpoints and guide: <https://github.com/moojink/openvla-oft/blob/main/LIBERO.md>

4 datasets (LIBERO-Spatial, LIBERO-Object, LIBERO-Goal, and LIBERO-10)

- Clone the OpenVLA and LIBERO repos.
- Run `./vla-scripts/finetune.py`
- Then run `./experiments/robot/libero/run_libero_eval.py`

The paper used 8 A100 or H100 GPUs with 80 GB memory. Takes about 1-2 days per dataset

Misha has fewer GPUs, so it's estimated to take about 7-8 days per dataset.

Qualitative Examples

ACT: Scoop raisins into bowl



Diffusion: Scoop pretzels into bowl



Qualitative Examples

After OFT training on OpenVLA

- L1 Regression is less noisy than diffusion, allows for more errors in training data



Autoregressive vs Diffusion

Aspect	Diffusion	Autoregressive
Examples	RDT-1B (1.2B), Diffusion Policy, π_0	RT-2, OpenVLA, π_0 -FAST
Prediction Method	Full trajectories via denoising	Token-by-token or action chunks
Action Distribution	Multimodal (multiple solutions)	Unimodal (single solution)
Inference Speed	Slow (10-100 steps, ~100-200ms)	Fast (single pass, < 10ms)
Best For	Dexterous tasks and contact rich control	Language-based tasks, fast deployment, general-purpose
Use when	Data-constrained	Compute-constrained

References

<https://openvla-oft.github.io/>

<https://blog.ml.cmu.edu/2025/09/22/diffusion-beats-autoregressive-in-data-constrained-settings/>

<https://github.com/moojink/openvla-oft/blob/main/LIBERO.md>

Pi0

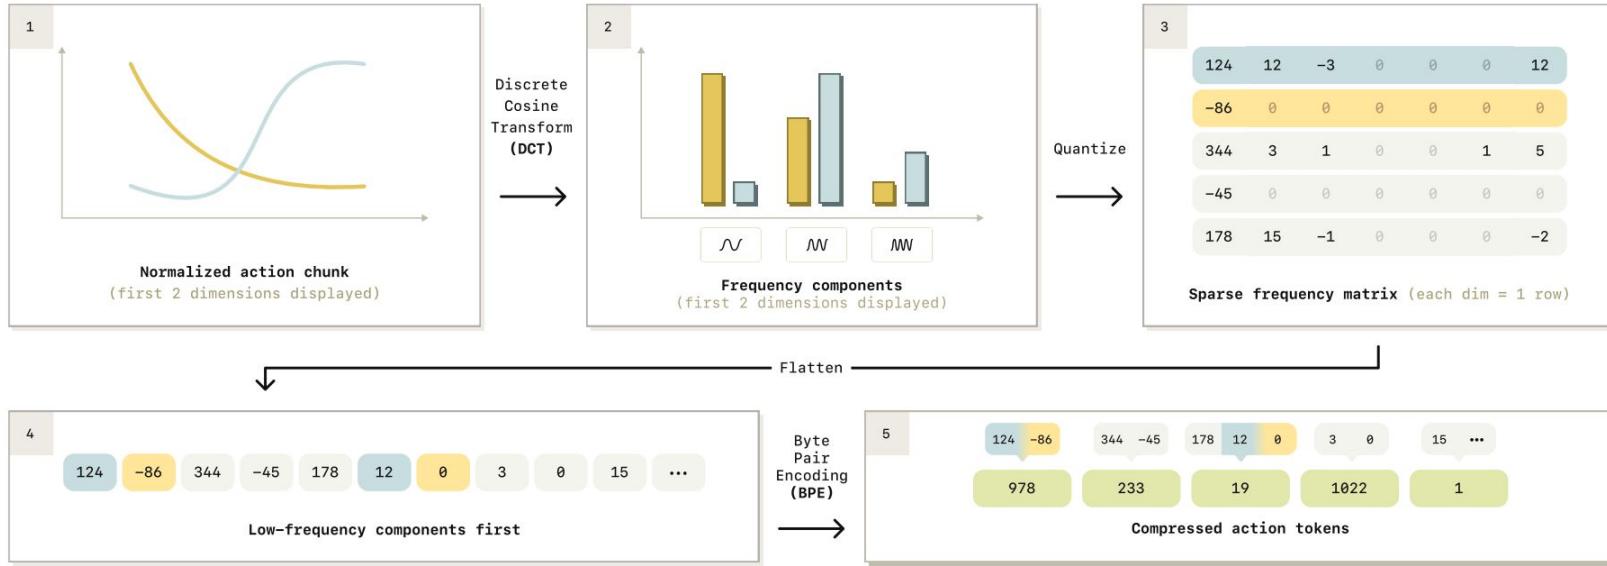
- Target generalist robotics policy
 - Data from many robot types is combined into the same model
- Key Components:
 - VLM backbone: PaliGemma (SigLIP + Gemma) + robotics inputs
 - Flow-matching action head allows for high frequency, continuous action
 - Two “experts”: robotics vs vision + language
- Emphasis on multimodal data labeling on whole trajectories and sub-trajectories for modality alignment

Pi0-FAST

- Diffusion / Flow matching training is very slow compared to autoregressive
- Autoregressive VLAs output discrete tokens that map to continuous actions
- Hard to capture high frequency behaviors
- Tackle with Discrete Cosine Transform (DCT) signal compression method
 - Variant of discrete Fourier Transform but only real (cosine) components, leading to better representation
- Leads to 5x faster training

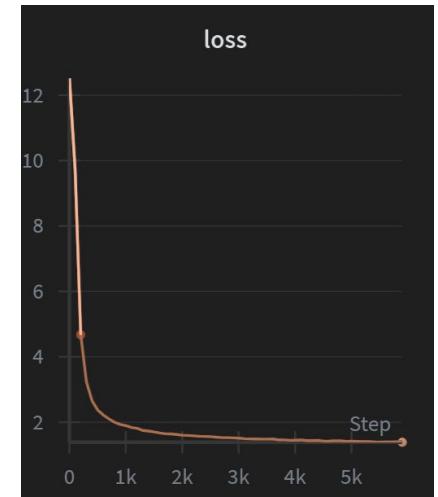
$$C(u) = a(u) \sum_{x=0}^{N-1} f(x) \cos\left[\frac{(2x+1)u\pi}{2N}\right]$$
$$u = 0, 1, \dots, N-1$$

$$a(u) = \begin{cases} \sqrt{\frac{1}{N}} & u = 0 \\ \sqrt{\frac{2}{N}} & u = 1, \dots, N-1 \end{cases}$$



Fine-tuning Pi0

- Physical Intelligence contains many checkpoints for Pi0, Pi0-Fast, Pi0.5 and finetuned versions on LIBERO, ALOHA, DROID, etc.
- Finetuning is very simple: convert dataset to LeRobot and create a finetuning config



Continuing from Pi0

- After Pi0, VLA robotic foundation models have built on their approach
- Nvidia GR00T N1
 - Diffusion Transformer Action Head
 - Auxiliary object detection loss to align vision modality
 - Action labels for human and synthetic video data generated by VQVAE
- Gemini Robotics
 - Similar to Pi0, VLA model with generalist capabilities (Zero-shot Transfer)
 - Uses a reasoning VLM to orchestrate and break down complex tasks

References

<https://arxiv.org/pdf/2410.24164v1>

<https://arxiv.org/pdf/2501.09747>

<https://tuul.ai/research/groot-n1-robotic-foundation-model>